




a Mount a curated ontology as retrieval geometry



b What each retriever keeps — and discards



c Structure-mapping retrieval with verifiable agency



d One universal loader; route across domains



Fig. 1 | A structure-mapping memory grounded in a curated ontology turns a generalist LLM into a verifiable specialist. SMA retrieves on the logical structure that vector RAG and knowledge graphs discard. **a**, a raw artifact is curated into an expert-maintained ontology — an is-a ascension lattice (geometry) with typed relations — and a case is encoded as functor statements placed on it, so a specific term can match a more general one. **b**, vector RAG collapses an item to one vector (discarding rarity and structure); a knowledge graph hops by adjacency (no subsumption); SMA keeps subsumption, higher-order structure, and rare-but-decisive detail. **c**, retrieval is a structural alignment: MAC returns an admissible-bound shortlist, FAC performs best-first structural alignment, and every hit yields a structural citation, with calibrated abstention and novelty flagging. **d**, one universal loader mounts any open ontology and a router selects the domain per query (HPO, GO, ATT&CK, CPC, FIBO, ChEBI).