



A golden-ontology structure-mapping memory turns a generalist LLM into a verifiable specialist



Three capabilities vector RAG and knowledge graphs structurally lack



Medicine LLM-QA, LLM & prompt fixed, memory swapped — SMA vs dense RAG: **accuracy 0.342 vs 0.100**

Generalizes across 5 domains: tail top-5 wins (med +0.33, gen +0.16, fin +0.17, cyber +0.07, legal +0.06; all Holm-sig.)